

# Performance Requirements

---

**Source File:** generated-documents\technical-design\performance-requirements.md

**Generated:** 16/07/2025 at 14:00:40

**Generated by:** Requirements Gathering Agent - PDF Converter

## PerformanceRequirements

---

**Generated by** adpa-enterprise-framework-automation v3.2.0

**Category:** technical-design

**Generated:** 2025-07-14T21:05:44.091Z

**Description:**

---

Certainly! Based on the provided project context and best practices for enterprise software performance requirements, here is a comprehensive set of **Performance Requirements** for the ADPA (Advanced Document Processing & Automation Framework):

---

## Performance Requirements: ADPA Enterprise Framework

---

### 1. Performance Goals

---

- **High Responsiveness:** Ensure all document generation, API, and integration operations meet enterprise-grade responsiveness

suitable for business-critical workflows.

- **Consistent Throughput:** Maintain steady throughput for concurrent document generation, AI requests, and integration tasks, even under peak usage.
  - **Enterprise Scalability:** Seamlessly scale horizontally to support large teams, multiple projects, and high-volume automation scenarios.
  - **Reliability Under Load:** Guarantee system stability and predictable degradation under high or unexpected load.
- 

## 2. Response Time Requirements

---

### 2.1 API Endpoints (Measured at 95th percentile)

- **Document Generation (POST /api/v1/generate):**
  - $\leq 3$  seconds for standard templates ( $\leq 10$  pages)
  - $\leq 8$  seconds for complex/AI-powered templates ( $\leq 30$  pages)
- **Templates Listing (GET /api/v1/templates):**
  - $\leq 500\text{ms}$
- **Confluence/SharePoint Publishing:**
  - $\leq 5$  seconds for documents  $\leq 5$  MB
- **Health/Readiness Endpoints:**
  - $\leq 200\text{ms}$
- **Admin Web Interface Page Loads:**
  - $\leq 1.0$  second for dashboard and core pages

### 2.2 CLI Operations

- **CLI Command Feedback:**
    - Initial response  $\leq 500\text{ms}$
    - Completion of standard document generation  $\leq 5$  seconds
- 

## 3. Throughput Expectations

---

- **Concurrent API Requests:**
    - Support at least 100 concurrent active API sessions without degradation.
  - **Document Generation:**
    - Minimum 20 documents/minute (with average document size  $\leq 10$  pages) per node.
  - **AI Provider Calls:**
    - At least 50 AI provider calls/minute per node, with automatic failover.
  - **Integration Operations:**
    - Simultaneous publishing to at least 3 integrations (e.g., Confluence, SharePoint, Adobe) without blocking.
- 

## 4. Scalability Requirements

---

- **Horizontal Scaling:**
    - The system must support horizontal scaling via stateless microservices.
    - Adding new API server/worker nodes should linearly increase throughput.
  - **Elastic Resource Management:**
    - Support dynamic scaling in cloud/containerized environments (e.g., Docker/Kubernetes).
  - **Multi-tenant Support:**
    - System must maintain isolation and fair resource allocation across multiple enterprise clients.
- 

## 5. Resource Utilization

---

- **CPU Utilization:**
  - Average CPU usage per node  $\leq 70\%$  under typical load; never exceeds 90% for  $>5$  minutes.
- **Memory Utilization:**

- Average memory usage  $\leq 75\%$  of available memory per service.
  - **Disk I/O:**
    - Document generation and storage processes should not exceed 80% of disk I/O bandwidth.
  - **Network:**
    - API and integration endpoints must not saturate available network bandwidth; utilize connection pooling and keep-alive.
- 

## 6. Load Handling

---

- **Graceful Degradation:**
    - Under overload, system must return 429 (Too Many Requests) or similar error, never crash.
  - **Backpressure:**
    - Implement request queueing and rate limiting (e.g., express-rate-limit) to avoid overload.
  - **Retry Strategy:**
    - Failed AI provider or integration calls must be retried with exponential backoff, up to 3 times.
- 

## 7. Caching Strategy

---

- **Template Caching:**
  - Frequently used document templates and assets must be cached in memory (e.g., Redis) for  $\leq 100\text{ms}$  retrieval.
- **AI Provider Response Caching:**
  - Cache non-personalized AI responses to reduce duplicate provider usage.
- **API Response Caching:**
  - Cache non-sensitive API responses (e.g., template lists) for at least 60 seconds.
- **Invalidation:**
  - Invalidate cache on template or configuration updates.

---

## 8. Performance Metrics

---

- **API Latency:**
    - 95th/99th percentile response times for all endpoints.
  - **Throughput:**
    - Requests per second (RPS) per node/service.
  - **Error Rate:**
    - $\leq 0.1\%$  failed requests under normal load.
  - **Resource Usage:**
    - CPU, memory, disk I/O, network per node.
  - **External API Latency:**
    - Average and max response time for each AI provider/integration.
  - **Queue Depth:**
    - Number of queued/rate-limited requests.
  - **Cache Hit Ratio:**
    - $\geq 90\%$  for template cache.
- 

## 9. Monitoring Requirements

---

- **Real-Time Dashboards:**
  - Integrate with monitoring tools (e.g., Prometheus, Grafana, Azure Monitor) for live visibility into all metrics above.
- **Health Checks:**
  - Expose `/api/v1/health` and `/api/v1/health/ready` endpoints, monitored every 30 seconds.
- **Alerting:**
  - Automated alerts for:
    - API latency > target thresholds (95th percentile)
    - Error rate > 0.5%
    - CPU or memory > 85% for > 10 minutes
    - Integration/AI provider failures
- **Log Aggregation:**

- Centralized logging (e.g., via Winston/Morgan/ELK) for all request/response cycles and error events.
  - **Audit Trails:**
    - Track document generation, publishing, and user actions for post-mortem analysis.
- 

## 10. Performance Testing Plan

---

- **Automated Baseline Tests:**
    - Run `npm run test:performance` after every release; validate all response time and throughput targets.
  - **Load Testing:**
    - Simulate 2x expected peak concurrent users (e.g., 200+ API sessions) for 1 hour; system must remain stable, with <1% error rate.
  - **Stress Testing:**
    - Incrementally increase load until resource limits are hit; system must degrade gracefully, not fail catastrophically.
  - **End-to-End Scenarios:**
    - Test complex workflows (multi-step document generation, integration publishing) under load.
  - **Integration Tests:**
    - Validate AI failover, caching effectiveness, and queueing under external provider/API slowness.
  - **Reporting:**
    - Generate detailed test reports with metrics, bottleneck identification, and recommendations.
- 

## Baseline Requirements Recap

---

- **API 95th percentile response time  $\leq 3s$  for standard ops**
- **Minimum 100 concurrent API sessions per node**
- **Average CPU  $\leq 70\%$ , Memory  $\leq 75\%$**
- **Graceful degradation with error messages under load**

- **Cache hit ratio  $\geq$  90% for templates**
  - **Automated, continuous performance testing and alerting**
- 

**Note:** All requirements must be reviewed and updated quarterly or upon significant architecture/infrastructure changes.

---

## **End of Performance Requirements**

---

Generated from generated-documents\technical-design\performance-requirements.md |  
Requirements Gathering Agent